

Problem 1: Python & Data Exploration

In [30]:

```
# Write and run your code here
import numpy as np
import matplotlib.pyplot as plt
iris = np.genfromtxt("data/iris.txt", delimiter=None) # Load the text file
Y = iris[:, -1] # target value is the last column
X = iris[:, 0:-1] # features are the other columns
```

In [31]:

```
# Part 1
X.shape
```

#We get 148 and 4 here. 148 means that there are 148 data points. 4 here means that there are 4 features each.

Out[31]:


(148, 4)

In [32]:

```
# Part 2
X1 = X[:, 0]
plt.hist(X1)
```

Out[32]:

```
(array([ 9., 20., 16., 27., 16., 25., 17., 10., 2., 6.]),
 array([4.3265579, 4.68918546, 5.05181302, 5.41444058, 5.77706814,
        6.1396957, 6.50232326, 6.86495082, 7.22757838, 7.59020594,
        7.9528335 ]),
 <BarContainer object of 10 artists>)
```




In [33]:

```
X2 = X[:, 1]
plt.hist(X2)
```

Out[33]:

```
(array([ 1., 10., 17., 28., 43., 23., 12., 9., 3., 2.]),
 array([2.0020454, 2.24764816, 2.49325092, 2.73885368, 2.98445644,
        3.2300592, 3.47566196, 3.72126472, 3.96686748, 4.21247024,
        4.458073 ]),
 <BarContainer object of 10 artists>)
```




In [34]:

```
X3 = X[:,2]
plt.hist(X3)
```

Out[34]:

```
(array([36., 13., 0., 3., 7., 25., 26., 21., 11., 6.]),
 array([1.0238316 , 1.61492121, 2.20601082, 2.79710043, 3.38819004,
        3.97927965, 4.57036926, 5.16145887, 5.75254848, 6.34363809,
        6.9347277 ]),
 <BarContainer object of 10 artists>)
```




In [35]:

```
X4 = X[:,3]
plt.hist(X4)
```

Out[35]:

```
(array([40., 8., 1., 7., 17., 23., 9., 20., 14., 9.]),
 array([0.12509037, 0.37029368, 0.615497 , 0.86070031, 1.10590362,
        1.35110693, 1.59631025, 1.84151356, 2.08671687, 2.33192019,
        2.5771235 ]),
 <BarContainer object of 10 artists>)
```



In [36]:

```
#Part 3
print(np.mean(X, axis = 0))
print(np.std(X, axis = 0))
```

```
[5.90010376 3.09893092 3.81955484 1.25255548]
[0.83340207 0.43629184 1.75405711 0.75877246]
```


In [37]:

```
#Part 4
def color_set():
    color = []
    for i in range(X.shape[0]):
        if Y[i] == 0 : color.append("g")
        elif Y[i] == 1: color.append("r")
        elif Y[i] == 2: color.append("b")
    return color

def Plot_graph(a,x,y):
    a.scatter(X[:, x - 1], X[:, y - 1], c = color)
    a.set_title("feature %d and %d" % (x, y))

def plot():
    fig, [a0,a1,a2] = plt.subplots(3,1)
    Plot_graph(a0, 1, 2)
    Plot_graph(a1, 1, 3)
    Plot_graph(a2, 1, 4)
    plt.show()

color = color_set()
plot()
```



Problem 2: k-Nearest Neighbor (kNN) exercise

The kNN classifier consists of two stages:

- During training, the classifier takes the training data and simply remembers it
- During testing, kNN classifies every test image by comparing to all training images and transferring the labels of the k most similar training examples
- The value of k is cross-validated

In this exercise you will implement these steps and understand the basic Image Classification pipeline, cross-validation, and gain proficiency in writing efficient, vectorized code.

In [38]:

```
# Run some setup code for this notebook.

import random
import numpy as np
from cs273p.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

# This is a bit of magic to make matplotlib figures appear inline in the notebook
# rather than in a new window.
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# Some more magic so that the notebook will reload external python modules;
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

In [39]:

```
%cd cs273p/datasets
!source get_datasets.sh
```

C:\Users\10630\273p\hw1\cs273p\datasets

'source' is not recognized as an internal or external command,
operable program or batch file.

In [40]:

```
%cd ../../..
```

```
C:\Users\10630\273p\hw1
```

In [41]:

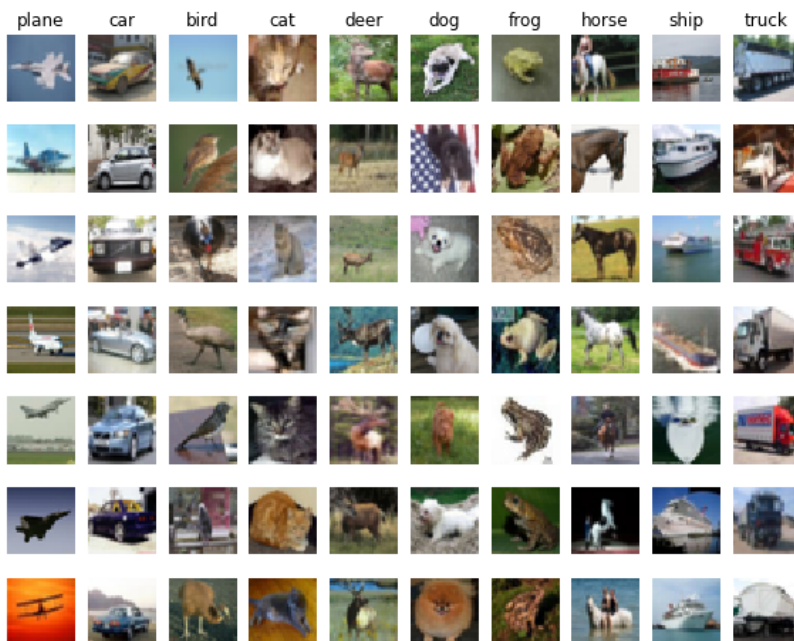
```
# Load the raw CIFAR-10 data.
cifar10_dir = './cs273p/datasets/cifar-10-batches-py'
X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

# As a sanity check, we print out the size of the training and test data.
print('Training data shape: ', X_train.shape)
print('Training labels shape: ', y_train.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

```
Training data shape: (50000, 32, 32, 3)
Training labels shape: (50000,)
Test data shape: (10000, 32, 32, 3)
Test labels shape: (10000,)
```

In [42]:

```
# Visualize some examples from the dataset.
# We show a few examples of training images from each class.
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
num_classes = len(classes)
samples_per_class = 7
for y, cls in enumerate(classes):
    idxs = np.flatnonzero(y_train == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(X_train[idx].astype('uint8'))
        plt.axis('off')
        if i == 0:
            plt.title(cls)
plt.show()
```



In [43]:

```
# Subsample the data for more efficient code execution in this exercise
num_training = 5000
mask = list(range(num_training))
X_train = X_train[mask]
y_train = y_train[mask]

num_test = 500
mask = list(range(num_test))
X_test = X_test[mask]
y_test = y_test[mask]
```

In [44]:

```
# Reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
print(X_train.shape, X_test.shape)
```

(5000, 3072) (500, 3072)

In [45]:

```
from cs273p.classifiers import KNearestNeighbor

# Create a kNN classifier instance.
# Remember that training a kNN classifier is a noop:
# the Classifier simply remembers the data and does no further processing
classifier = KNearestNeighbor()
classifier.train(X_train, y_train)
```

We would now like to classify the test data with the kNN classifier. Recall that we can break down this process into two steps:

1. First we must compute the distances between all test examples and all train examples.
2. Given these distances, for each test example we find the k nearest examples and have them vote for the label

Lets begin with computing the distance matrix between all training and test examples. For example, if there are **Ntr** training examples and **Nte** test examples, this stage should result in a **Nte x Ntr** matrix where each element (i,j) is the distance between the i-th test and j-th train example.

First, open `cs273p/classifiers/k_nearest_neighbor.py` and implement the function `compute_distances_two_loops` that uses a (very inefficient) double loop over all pairs of (test, train) examples and computes the distance matrix one element at a time.

In [46]:


```
# Open cs273p/classifiers/k_nearest_neighbor.py and implement
# compute_distances_two_loops.

# Test your implementation:
dists = classifier.compute_distances_two_loops(X_test)
print(dists.shape)
```

(500, 5000)

In [47]:

```
# We can visualize the distance matrix: each row is a single test example and
# its distances to training examples
plt.imshow(dists, interpolation='none')
plt.show()
```



Inline Question #1: Notice the structured patterns in the distance matrix, where some rows or columns are visible brighter. (Note that with the default color scheme black indicates low distances while white indicates high distances.)

- What in the data is the cause behind the distinctly bright rows?
- What causes the columns?

Your Answer: 1). The cause of it is the difference between distance. The rows that are brighter means that the distance there is bigger. 2). For this part, the distance between train data and test data here is too big, so the graph here is mostly dark.

In [48]:

```
# Now implement the function predict_labels and run the code below:
# We use k = 1 (which is Nearest Neighbor).
y_test_pred = classifier.predict_labels(dists, k=1)

# Compute and print the fraction of correctly predicted examples
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))
```

Got 137 / 500 correct => accuracy: 0.274000

You should expect to see approximately 27% accuracy. Now lets try out a larger k, say k = 5 :

In [49]:

```
y_test_pred = classifier.predict_labels(dists, k=5)
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))
```

Got 139 / 500 correct => accuracy: 0.278000

You should expect to see a slightly better performance than with $k = 1$.

In [50]:

```
# Now Lets speed up distance matrix computation by using partial vectorization
# with one loop. Implement the function compute_distances_one_loop and run the
# code below:
dists_one = classifier.compute_distances_one_loop(X_test)

# To ensure that our vectorized implementation is correct, we make sure that it
# agrees with the naive implementation. There are many ways to decide whether
# two matrices are similar; one of the simplest is the Frobenius norm. In case
# you haven't seen it before, the Frobenius norm of two matrices is the square
# root of the squared sum of differences of all elements; in other words, reshape
# the matrices into vectors and compute the Euclidean distance between them.
difference = np.linalg.norm(dists - dists_one, ord='fro')
print('Difference was: %f' % (difference, ))
if difference < 0.001:
    print('Good! The distance matrices are the same')
else:
    print('Uh-oh! The distance matrices are different')
```

Difference was: 0.000000

Good! The distance matrices are the same

In [51]:

```
# Now implement the fully vectorized version inside compute_distances_no_loops
# and run the code
dists_two = classifier.compute_distances_no_loops(X_test)

# check that the distance matrix agrees with the one we computed before:
difference = np.linalg.norm(dists - dists_two, ord='fro')
print('Difference was: %f' % (difference, ))
if difference < 0.001:
    print('Good! The distance matrices are the same')
else:
    print('Uh-oh! The distance matrices are different')
```

Difference was: 0.000000

Good! The distance matrices are the same

In [52]:

```
# Let's compare how fast the implementations are
def time_function(f, *args):
    """
    Call a function f with args and return the time (in seconds) that it took to execute.
    """
    import time
    tic = time.time()
    f(*args)
    toc = time.time()
    return toc - tic

two_loop_time = time_function(classifier.compute_distances_two_loops, X_test)
print('Two loop version took %f seconds' % two_loop_time)

one_loop_time = time_function(classifier.compute_distances_one_loop, X_test)
print('One loop version took %f seconds' % one_loop_time)

no_loop_time = time_function(classifier.compute_distances_no_loops, X_test)
print('No loop version took %f seconds' % no_loop_time)

# you should see significantly faster performance with the fully vectorized implementation
```

Two loop version took 20.480723 seconds

One loop version took 59.934474 seconds

No loop version took 0.259705 seconds

Cross-validation

We have implemented the k-Nearest Neighbor classifier but we set the value $k = 5$ arbitrarily. We will now determine the best value of this hyperparameter with

cross-validation.

In [53]:

```

num_folds = 5
k_choices = [1, 3, 5, 8, 10, 12, 15, 20, 50, 100]

X_train_folds = []
y_train_folds = []
#####
# TODO:
# Split up the training data into folds. After splitting, X_train_folds and
# y_train_folds should each be lists of length num_folds, where
# y_train_folds[i] is the label vector for the points in X_train_folds[i].
# Hint: Look up the numpy array_split function.
#####
X_train_folds=np.array_split(X_train, 5)
y_train_folds=np.array_split(y_train, 5)
#####
#                               END OF YOUR CODE
#####

# A dictionary holding the accuracies for different values of k that we find
# when running cross-validation. After running cross-validation,
# k_to_accuracies[k] should be a list of length num_folds giving the different
# accuracy values that we found when using that value of k.
k_to_accuracies = {}

#####
# TODO:
# Perform k-fold cross validation to find the best value of k. For each
# possible value of k, run the k-nearest-neighbor algorithm num_folds times,
# where in each case you use all but one of the folds as training data and the
# last fold as a validation set. Store the accuracies for all fold and all
# values of k in the k_to_accuracies dictionary.
#####
for k in k_choices:
    for f in range(num_folds):
        Xhold = X_train_folds[f]
        Yhold = y_train_folds[f]
        tempX = np.concatenate(X_train_folds[:f] + X_train_folds[f + 1:])
        tempY = np.concatenate(y_train_folds[:f] + y_train_folds[f + 1:])
        classifier = KNearestNeighbor()
        classifier.train( tempX, tempY )
        distancee = classifier.compute_distances_no_loops(Xhold)
        predistance = classifier.predict_labels(distancee, k=k)
        exph = Xhold.shape[0]
        accuracy = float( np.sum(predistance == Yhold)) / exph
        k_to_accuracies.setdefault(k, []).append(accuracy)
#####
#                               END OF YOUR CODE
#####

# Print out the computed accuracies
for k in sorted(k_to_accuracies):
    for accuracy in k_to_accuracies[k]:
        print('k = %d, accuracy = %f' % (k, accuracy))

```




```
k = 1, accuracy = 0.263000
k = 1, accuracy = 0.257000
k = 1, accuracy = 0.264000
k = 1, accuracy = 0.278000
k = 1, accuracy = 0.266000
k = 3, accuracy = 0.239000
k = 3, accuracy = 0.249000
k = 3, accuracy = 0.240000
k = 3, accuracy = 0.266000
k = 3, accuracy = 0.254000
k = 5, accuracy = 0.248000
k = 5, accuracy = 0.266000
k = 5, accuracy = 0.280000
k = 5, accuracy = 0.292000
k = 5, accuracy = 0.280000
k = 8, accuracy = 0.262000
k = 8, accuracy = 0.282000
k = 8, accuracy = 0.273000
k = 8, accuracy = 0.290000
k = 8, accuracy = 0.273000
k = 10, accuracy = 0.265000
k = 10, accuracy = 0.296000
k = 10, accuracy = 0.276000
k = 10, accuracy = 0.284000
k = 10, accuracy = 0.280000
k = 12, accuracy = 0.260000
k = 12, accuracy = 0.295000
k = 12, accuracy = 0.279000
k = 12, accuracy = 0.283000
k = 12, accuracy = 0.280000
k = 15, accuracy = 0.252000
k = 15, accuracy = 0.289000
k = 15, accuracy = 0.278000
k = 15, accuracy = 0.282000
k = 15, accuracy = 0.274000
k = 20, accuracy = 0.270000
k = 20, accuracy = 0.279000
k = 20, accuracy = 0.279000
k = 20, accuracy = 0.282000
k = 20, accuracy = 0.285000
k = 50, accuracy = 0.271000
k = 50, accuracy = 0.288000
k = 50, accuracy = 0.278000
k = 50, accuracy = 0.269000
k = 50, accuracy = 0.266000
k = 100, accuracy = 0.256000
k = 100, accuracy = 0.270000
k = 100, accuracy = 0.263000
k = 100, accuracy = 0.256000
k = 100, accuracy = 0.263000
```

In [54]:

```
# plot the raw observations
for k in k_choices:
    accuracies = k_to_accuracies[k]
    plt.scatter([k] * len(accuracies), accuracies)

# plot the trend line with error bars that correspond to standard deviation
accuracies_mean = np.array([np.mean(v) for k,v in sorted(k_to_accuracies.items())])
accuracies_std = np.array([np.std(v) for k,v in sorted(k_to_accuracies.items())])
plt.errorbar(k_choices, accuracies_mean, yerr=accuracies_std)
plt.title('Cross-validation on k')
plt.xlabel('k')
plt.ylabel('Cross-validation accuracy')
plt.show()
```



In [55]:

```
# Based on the cross-validation results above, choose the best value for k,
# retrain the classifier using all the training data, and test it on the test
# data. You should be able to get above 28% accuracy on the test data.
best_k = k_choices[np.argmax(accuracies_mean)]

classifier = KNearestNeighbor()
classifier.train(X_train, y_train)
y_test_pred = classifier.predict(X_test, k=best_k)

# Compute and display the accuracy
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))
```

Got 141 / 500 correct => accuracy: 0.282000

Problem 3: Naïve Bayes Classifiers

You don't need to code this question. You can either type your answer or attach an image of hand written solution here.

In [56]:

```
#Problem 1
#P(y = 1) = 4/10
#P(y = -1) = 6/10
#P(x1 = 0 | y = 1) = 1/4
#P(x1 = 0 | y = -1) = 2/4
#P(x1 = 1 | y = 1) = 3/4
#P(x1 = 1 | y = -1) = 1/2
#P(x2 = 0 | y = 1) = 1
#P(x2 = 0 | y = -1) = 1/6
#P(x2 = 1 | y = 1) = 0/4
#P(x2 = 1 | y = -1) = 5/6
#P(x3 = 0 | y = 1) = 1/4
#P(x3 = 0 | y = -1) = 2/6
#P(x3 = 1 | y = 1) = 3/4
#P(x3 = 1 | y = -1) = 4/6
#P(x4 = 0 | y = 1) = 2/4
#P(x4 = 0 | y = -1) = 1/6
#P(x4 = 1 | y = 1) = 2/4
#P(x4 = 1 | y = -1) = 5/6
#P(x5 = 0 | y = 1) = 3/4
#P(x5 = 0 | y = -1) = 4/6
#P(x5 = 1 | y = 1) = 1/4
#P(x5 = 1 | y = -1) = 2/6
```

In [57]:

```
#problem 2
#First part (0 0 0 0 0)
#P(y = 1, x = (0 0 0 0 0))
# = P(y=1) * P(x1 = 0 | y = 1) * P(x2 = 0 | y = 1) * P(x3 = 0 | y = 1) * P(x4 = 0 | y = 1) * P(x5 = 0 | y = 1) = 0.0094
#P(y = -1, x = (0 0 0 0 0))
# = P(y= -1) * P(x1 = 0 | y = -1) * P(x2 = 0 | y = -1) * P(x3 = 0 | y = -1) * P(x4 = 0 | y = -1) * P(x5 = 0 | y = -1) = 0.0019
#0.0094 > 0.0019, so it belongs to class 1
#Second part (1 1 0 1 0)
#P(y = 1, x = (1 1 0 1 0)) = 0
#P(y = -1, x = (1 1 0 1 0)) = 0.46
# 0 < 0.46, so it belongs to class -1
```

In [58]:

```
#Problem 3
# P(y = 1 | x = (1 1 0 1 0)) = 0 (we have calculated that P(y = 1, x = (1 1 0 1 0)) = 0, so this is 0)
#The answer is 0
```

In []:

```
#Problem 4
# The reason is that the information we need to store is too much for this situation.
# This also means that we need to calculate many probability combinations.
# What is more, most of the feature are independent, so Naive Bayes would be much better here.
```

In []:

```
#Problem 5
# The answer is we should not do it. For the reason that the features in Naive Bayes Classifier are independent.
# It means that we do not need all the features to predict.
# This means that the retrain here is not necessary at all.
```

In [1]:

```
# Statement of Collaboration
# Name I had discussion with Yanghan-Deng, We dicussed about the function we should use for two Loops and also no loop above.
#Yet we did not talk about the answer with each other.
```